

TRANSFER LEARNING

Dr. Brian Mc Ginley



INTRODUCTION TO TRANSFER LEARNING

- Transfer learning is a machine learning technique where a model trained on one task is reused as a starting point for a model on a different task.
- Instead of starting the learning process from scratch, transfer learning allows us to leverage knowledge gained from solving one problem to solve a different but related problem.
- Humans use transferrable skills all the time



WHY TRANSFER LEARNING?

- **Data Efficiency:** Transfer learning enables effective learning with smaller datasets by leveraging knowledge from larger datasets.
- **Faster Training:** Starting with pre-trained weights reduces training time significantly.
- **Improved Performance:** Transfer learning often leads to better performance, especially in domains where labelled data is scarce.



HOW TRANSFER LEARNING WORKS

- **Step 1:** Pre-train a model on a large dataset with a source task.
- **Step 2:** Fine-tune the pre-trained model on a smaller dataset with a target task.
- **Step 3:** Optionally, adapt the model architecture or add new layers for the target task.




ADDITIONAL RESOURCES

- Keras Documentation:
 - https://keras.io/guides/transfer_learning/
- TensorFlow Tutorials:
 - https://www.tensorflow.org/tutorials/images/transfer_learning



STEPS



OBTAINING THE MODEL

- Keras pre-trained models:
- There are more than two dozen pre-trained models available from Keras.
 - <https://keras.io/api/applications/>

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2
ResNet101V2	171	77.2%	93.8%	44.7M	205	72.7	5.4
ResNet152	232	76.6%	93.1%	60.4M	311	127.4	6.5
ResNet152V2	232	78.0%	94.2%	60.4M	307	107.5	6.6
InceptionV3	92	77.9%	93.7%	23.9M	189	42.2	6.9
InceptionResNetV2	215	80.3%	95.3%	55.9M	449	130.2	10.0
MobileNet	16	70.4%	89.5%	4.3M	55	22.6	3.4
MobileNetV2	14	71.3%	90.1%	3.5M	105	25.9	3.8
DenseNet121	33	75.0%	92.3%	8.1M	242	77.1	5.4
DenseNet169	57	76.2%	93.2%	14.3M	338	96.4	6.3



OBTAINING THE MODEL

- Keras pre-trained models:
- There are more than two dozen pre-trained models available from Keras.
 - <https://keras.io/api/applications/>
- Mobile Net Example
 - That's with default values. include_top means: include the fully-connected layers - often we want this as False.
 - Weights are the initial weights. If you do not use imagenet, then it will use random weights and you will need to train all these layers.

```
model = tf.keras.applications.MobileNet(  
    input_shape=None,  
    alpha=1.0,  
    depth_multiplier=1,  
    dropout=0.001,  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation="softmax",  
)
```



OBTAINING THE MODEL

- You can also use models from TensorFlow Hub - <https://www.tensorflow.org/hub>

```
model = tf.keras.Sequential([
    hub.KerasLayer("https://tfhub.dev/google/imagenet/mobilenet_v2_100_224/feature_vector",
                   trainable=False),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```



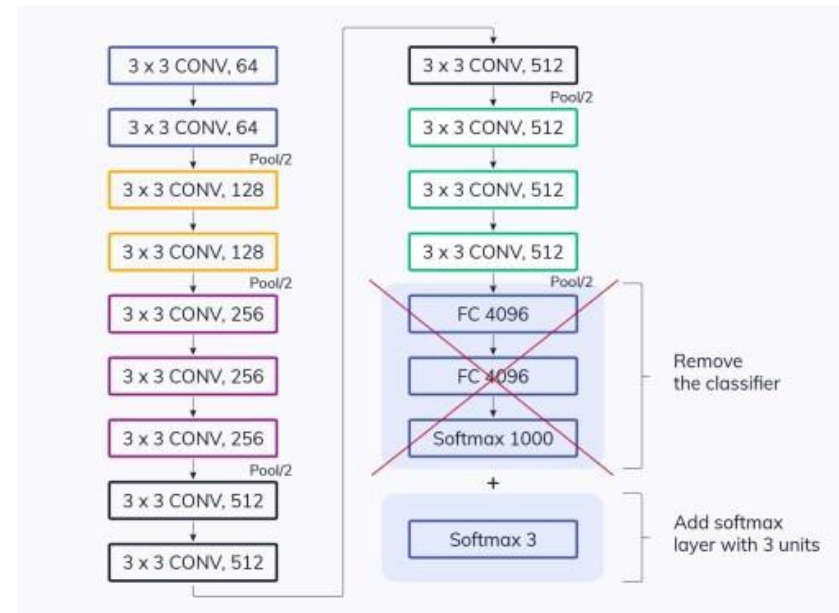
OBTAINING THE MODEL

- You can also use models from TensorFlow Hub - <https://www.tensorflow.org/hub>
- There are pretrained word-embeddings available, e.g. Google's word2vec
 - <https://code.google.com/archive/p/word2vec/>
- And for text-processing tasks there are pretrained models at
 - <https://github.com/huggingface>
 - Stanza another option <https://github.com/stanfordnlp/stanza/>
- Lots of pre-trained models also at: <https://www.kaggle.com/models>
 - Inception
 - VGG
 - ResNet
 - EfficientNet
 - MobileNet
 - YOLO etc.



CREATE BASE MODEL

- You instantiate the base model with pre-trained weights.
- The base model will usually have more units in the final output layer than you require so you have to at least remove the final output layer.
- Later on, you will add a final output layer that is compatible with your problem.



FREEZING LAYERS

- You need to freeze the layers from the pre-trained.
- If you do not, the weights will be retrained which defeats the whole purpose of transfer learning!

```
base_model.trainable = False
```

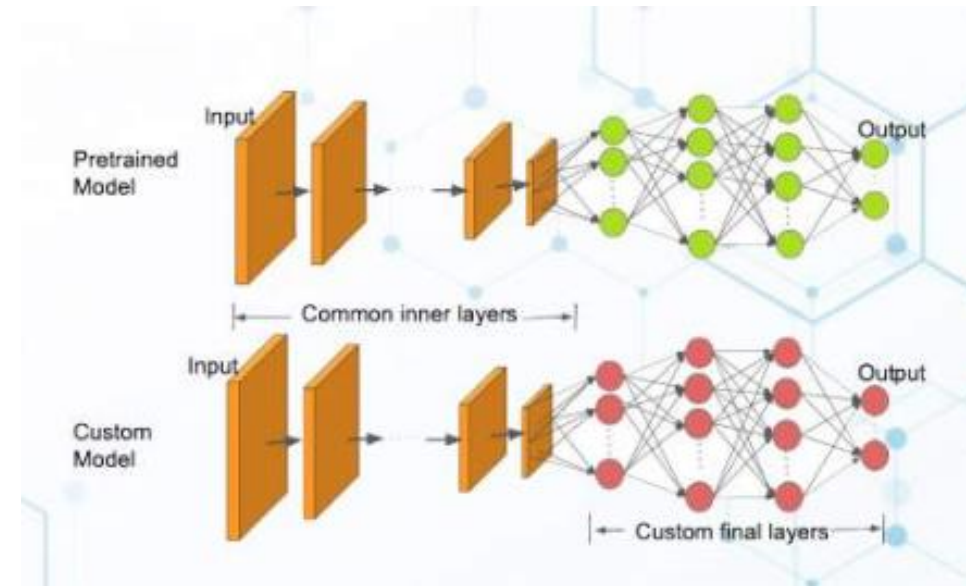
- You can also do it individually by layer(s)

```
for layer in base_model.layers:  
    layer.trainable = False
```



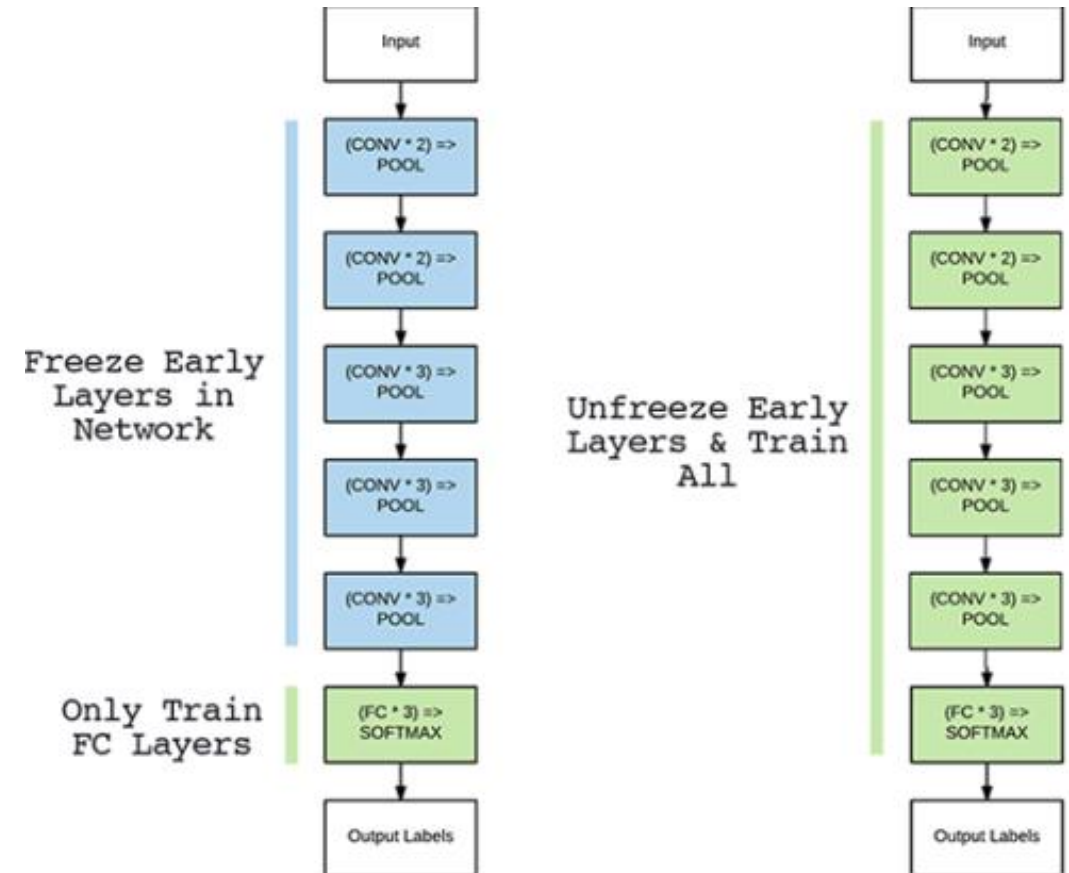
TRAIN NEW LAYERS

- You will add some new dense layers (another hyperparameter)
- Most importantly, a final dense layer with units corresponding to the number of outputs expected by your model.
- These new layer weights are the ones you will be training.
- Make sure the output from your base model will work as the input for your new layers.
- Compile and train, you now will have a model that can make predictions on your dataset.



FINE TUNING

- Optionally, you can improve its performance through fine-tuning.
- Fine-tuning is done by unfreezing the base model or part of it and training the entire model again on the whole dataset at a very low learning rate.



IMPORTANT THINGS TO THINK OF

- Always pay close attention to the base model's structure and know what it is supposed to do.
- Make sure the inputs you will be sending to the base model match what is expected.
- Are the image sizes 224x224 ? Or something else? Or is the model flexible enough to do something else?
- Does the model expect data to be normalised to $[0,1]$ or $[-1,1]$ or something else. Often the keras application will have a preprocess input function so read the manual.



IMPORTANT THINGS TO THINK OF

- Data Augmentation often is helpful. You may want to do this to your dataset, or you can build it into your model as before.

```
data_augmentation = tf.keras.Sequential([  
    keras.layers.preprocessing.RandomFlip("horizontal"),  
    keras.layers.preprocessing.RandomRotation(0.1),  
])
```



EXAMPLE

- This is taken from https://keras.io/guides/transfer_learning/, using the Xception model and some of my own alterations that I thought may help
- Instantiate the base model

```
base_model = tf.keras.applications.Xception(  
weights='imagenet',  
input_shape=(150, 150, 3),  
include_top=False) # Do not include the ImageNet classifier at the top
```

- Freeze the base model

```
base_model.trainable = False
```



EXAMPLE

- Create a new model on top, first let's set the size of the input images

```
inputs = tf.keras.Input(shape=(150, 150, 3))
```

- If you want to include data augmentation as part of the model you can do it here, https://www.tensorflow.org/guide/keras/transfer_learning includes it as part of the model while the tutorial I'm following does not.

```
x = data_augmentation(inputs)
```



EXAMPLE

- Make sure the data has been normalised as expected $[-1,1]$ or $[0,1]$ etc, use the function to aid you

```
# If you did not have the above data augmentation step
x = tf.keras.applications.xception.preprocess_input(inputs)
# If you did
x = tf.keras.applications.xception.preprocess_input(x)
```

- Alternatively, you can use a rescaling layer, for Xception

```
scale_layer = keras.layers.Rescaling(scale=1/127.5, offset=1)
x = scale_layer(inputs)
```



EXAMPLE

- Ensure that the base model is running in inference mode so that batch normalization layers are not updated during the fine-tuning stage (set 'training=False')
- Convert features from the base model to vectors, using 'GlobalAveragePooling2D' (a Flatten layer may work here too)
- Add whatever Dense layers, most importantly the final dense layer - this can have a softmax activation function
- We could add Dropout here too
- Finally put the model together



EXAMPLE

```
# We make sure that the base model is running in inference mode here,  
# by passing `training=False`. This is important for fine-tuning  
x = base_model(x, training=False)  
  
# Convert features of shape `base_model.output_shape[1:]` to vectors  
x = tf.keras.layers.GlobalAveragePooling2D()(x)  
  
# A Dense classifier with 10 units  
outputs = tf.keras.layers.Dense(10, activation='softmax')(x)  
  
model = tf.keras.Model(inputs, outputs)
```



EXAMPLE

- The previous needs to compile and fit (as done previously). Now after doing that we want to fine-tune the model we then do:

```
base_model.trainable = True
```

- Sets base model now to trainable, that's all layers. You could do individual layers by

```
freeze_layers_up_to = 10
for layer in base_model.layers[freeze_layers_up_to:]:
    layer.trainable = True
```

- You need to compile the model again, this time with a low learning rate

```
model.compile(
    optimizer=keras.optimizers.Adam(1e-5), # Low learning rate
    loss=keras.losses.CategoricalCrossentropy(from_logits=False),
    metrics=['accuracy'],
)
```

